

meeting-minutes skill: evaluation and benchmark results

3 test cases, each run with and without the skill, on Sonnet 5.5. Graded against 12 checks. Skill version 5c0d71c1 · 2026-10-09

Evaluated with Skills 2.0 - Anthropic skill-creator's eval and benchmark tooling: evals.json test cases, parallel isolated runs, grader agent, blind comparator agent (A/B), analyzer agent, aggregate_benchmark, run_eval trigger tests and the eval viewer. The step-by-step record is on page 2.

Skill under test	meeting-minutes
Description	Turns a meeting transcript or rough notes into minutes.md in the team's standard format - attendees, decisions, action items with a named owner and an absolute due date, and open questions. Use this whenever someone asks for minutes, a meeting summary, MoM, action items or follow-ups from a call, transcript or notes. Not for drafting agendas, emails to attendees or summarising documents that are not meetings.
Description check	412/1024 characters · says when to use: yes · says when not to use: yes · review: too narrow
Size	SKILL.md 35 lines (~369 tokens loaded when triggered); bundled files: none
Test cases	3 cases from evals/evals.json

KEEP	Is the skill still useful? Still useful. On Sonnet 5.5 the skill raises the pass rate by +25 pts; 3 checks pass only with the skill. 8 checks pass with or without it.
------	--

Model	With skill	Without skill	Uplift	Checks only the skill gets right	Checks the model passes unaided	Checks worse with skill	Token overhead	Verdict
Sonnet 5.5	92%	67%	+25 pts	3 (behavior 1, org knowledge 1, quality 1)	8	0	+3%	KEEP

A/B benchmark: blind comparison

For each case and run, Sonnet 5.5 saw the with-skill and the without skill output as A and B in random order (skill-creator's blind comparator) and picked the better one. Win rate counts ties as half. Skills 2.0 reading: 70% or more keep, 50-70% refine or question it, under 50% remove or rewrite.

Model	Comparisons	Skill wins	Ties	Baseline wins	Win rate	Reading
Sonnet 5.5	2	2	0	0	100%	KEEP

Skills 2.0 evaluation steps, as run

The Build, Eval, A/B benchmark and trigger-optimization workflow from the Skills 2.0 guide, with the skill-creator component behind each step and what it produced.

Stage	Step	Skills 2.0 component	Result	Status
Build	Validate the skill	scripts/quick_validate.py	SKILL.md valid; description 412/1024 chars	done
Eval	1. Define test cases	evals/evals.json (skill-creator schema)	3 cases, 12 checks; types: edge, should_not_trigger, standard; source: evals/evals.json - Skills 2.0 guidance: 5-10 cases including edge cases	warn
Eval	2. Run the eval (parallel, clean contexts)	headless Claude Code, one isolated agent per run	6 runs on Sonnet 5.5; up to 4 in parallel, each in an empty folder	done
Eval	Grade outputs	agents/grader.md + check scripts	pass rate Sonnet 5.5 92%; 1 failed check(s) with the skill	done
Eval	Metrics: response time and tokens per case	scripts/aggregate_benchmark.py -> benchmark.json	985,099 tokens metered in total, \$0.76	done
Eval	3. Fix failures	analyzer agent (agents/analyzer.md)	5 fix suggestion(s), 3 weak spot(s) listed in the report	done
Eval	4. Re-run until the threshold is met	target for a preference skill: 80%	Sonnet 5.5 meets (92%)	done
A/B benchmark	Blind comparison vs without skill	agents/comparator.md (blind A/B)	win rate Sonnet 5.5 100% (KEEP)	done
A/B benchmark	Interpret: keep, refine or remove	usefulness verdict (checks) + win rate	KEEP	done
Trigger optimization	Test borderline prompts	scripts/run_eval.py (trigger evals)	6/12 correct, recall 0%, false triggers 0%	done
Trigger optimization	Diagnose and revise the description	description review (analyzer)	too narrow; rewrite suggested	done
Review	Inspect every output	eval-viewer/generate_review.py	review.html in the run artifact	done
Maintain	Version in git with results	human gate -> eval-results/, history.json, registry	skill version 5c0d71c15485; stored after approval	pending
Maintain	Re-benchmark	daily scheduler (due / new model) + PR trigger	next due 2026-11-08	scheduled

Headline

On Sonnet 5.5 the skill lifted the format-sensitive checks: it produced the Confidential line, Open questions section and house layout that the baseline missed in case 1. The main caveat is trigger reliability: the skill auto-triggered on 0 of 6 genuine minutes requests in the trigger eval (recall 0%), so in real use it will rarely load unprompted.

Model	Config	Checks passed	Pass rate	Cases fully passed	Total tokens	Total seconds
Sonnet 5.5	With skill	11 / 12	92%	2 / 3	144,429	29.1
Sonnet 5.5	Without skill	8 / 12	67%	1 / 3	140,416	29.4
Sonnet 5.5	Uplift	+3	+25 pts	+1	+3% tokens	-1% time

Where the skill made the difference

- Case 1: with the skill, minutes.md has the Confidential - internal distribution only line because Victaulic's Tom is named; the baseline omitted it.
- Case 1: the skill produced a literal 'Open questions' section containing the dbt 1.9 upgrade and the Power BI request; the baseline used 'Open items' and wrote 'Decision: none today' inside the dbt section.
- Case 1: the skill's output uses the fixed Decisions / Action items / Open questions headings and ISO due dates (2026-10-09, 2026-10-16). The baseline wrote 'Fri 9 Oct' and 'End of next week' in the table.
- Case 1: the unowned dbt macro check is marked UNASSIGNED with an open-question entry. The baseline wrote 'Unassigned' and 'Not named (Rahul presumably)', which guesses an owner.
- Case 2: both runs passed everything; the unowned calendar-invite action was handled correctly even without the skill, so this edge case does not discriminate.
- Case 3: the with-skill run did not invoke the skill and produced a clean agenda. The baseline agenda included decision/action content that tripped the no-action-table check.

Where it does not help

- Case 1 check 1 failed with the skill: the validator rejected the UNASSIGNED row with a TBD due date, though the skill itself prescribes UNASSIGNED with no date. Fix: let the checker accept UNASSIGNED/TBD, or have the skill state the exact placeholder (for example 'UNASSIGNED | none').
- Case 2 gave no uplift (3/3 in both configs). Fix: add a harder edge case, such as two owners on one action, a date with no month, or 'next Tuesday' relative dates.
- Trigger eval: recall was 0 of 6 genuine requests, so the skill's gains only appear when it is forced to load. Fix the description (see description_review).

What still needs the skill, and what models now do unaided

A check counts for a model when it passes in at least half the runs. **Only with skill** is the value the skill still adds; **unaided** means the model passes it without the skill, so that guidance can be trimmed.

Case	Check	Category	Sonnet 5.5
1	A 'Confidential - internal distribution only' line appears because Victaulic is named	org knowledge	Only with skill
1	The dbt 1.9 upgrade and Power BI support appear under Open questions, not under Decisions	quality	Only with skill
3	No minutes.md, decisions or action-item table is produced	behavior	Only with skill
1	minutes.md has Decisions, Action items and Open questions, and every action item has one owner and a YYYY-MM-DD due date	org knowledge	Fails both
1	No small talk (coffee joke, audio check) is included	quality	Unaided
1	Rahul's hourly-refresh action is due 2026-10-09 and Meera's Power BI estimate is due 2026-10-16	org knowledge	Unaided
1	Retiring the old S3 bucket on 2026-10-31 is recorded as a decision	quality	Unaided
2	Moving standup to 10am is recorded as a decision	quality	Unaided
2	Priya's Q3 numbers action is due 2026-10-31	quality	Unaided
2	The calendar-invite action has owner UNASSIGNED and is also listed under Open questions	org knowledge	Unaided
3	Skill not invoked (automatic check)	behavior	Unaided
3	The reply is an agenda with the three topics and time allocations adding up to about 45 minutes	quality	Unaided

Side by side, per case

Case	Description	Checks	Sonnet 5.5 with	Sonnet 5.5 without	Sonnet 5.5 uplift
1	Minutes from a messy transcript	6	5/6	3/6	+2
2	Action nobody took	3	3/3	3/3	+0
3	Draft an agenda (should not trigger)	3	3/3	2/3	+1

Tokens and time per case

Case	Sonnet 5.5 tokens with / without	Sonnet 5.5 seconds with / without
1	54,943 / 53,407	10.9 / 11.6
2	53,032 / 34,214	6.7 / 5.0
3	36,454 / 52,795	11.5 / 12.8

Token usage

Model	Config	Runs	Input	Output	Cache read	Cache write	Total tokens	Cost
Sonnet 5.5	With skill	3	16	3,625	123,379	17,409	144,429	\$0.104
Sonnet 5.5	Without skill	3	16	3,462	121,477	15,461	140,416	\$0.098

Stage	Calls	Input	Output	Cache read + write	Total tokens	Cost
Skill runs (executor)	6	32	7,087	277,726	284,845	\$0.202
Grading (grader agent)	6	40	10,908	451,649	462,597	\$0.363
Blind A/B comparison (comparator agent)	2	12	3,085	127,201	130,298	\$0.102
Analysis (analyzer agent)	1	8	3,220	104,131	107,359	\$0.094
Total	15	92	24,300	960,707	985,099	\$0.761

Not metered: Trigger evals (run_eval.py streams events and stops early; usage is not reported).

What differed in each case

Case	Description	Notes
1	Minutes from a messy transcript	With skill: Confidential line, Open questions section, ISO dates and UNASSIGNED owner; the only failed check is the validator rejecting 'TBD' on the unowned row. Without: no confidentiality line, 'Open items' instead of Open questions, non-ISO dates and a guessed owner for the S3 task.
2	Action nobody took	With skill: decision recorded, Priya due 2026-10-31, calendar invite UNASSIGNED and listed in Open questions. Without: the same content in a differently named file; all checks pass in both.
3	Draft an agenda (should not trigger)	With skill: the skill stayed unused and a time-boxed agenda was written. Without: a valid agenda that also failed the 'no decisions or action-item table' check.

Check-by-check results

PASS/FAIL for each check on every model. Category marks where the value comes from (asset, org knowledge, safety, behavior, quality).

Case	Check	Category	Sonnet 5.5 with	Sonnet 5.5 without
1	minutes.md has Decisions, Action items and Open questions, and every action item has one owner and a YYYY-MM-DD due date [script]	org knowledge	FAIL	FAIL
1	Rahul's hourly-refresh action is due 2026-10-09 and Meera's Power BI estimate is due 2026-10-16	org knowledge	PASS	PASS
1	The dbt 1.9 upgrade and Power BI support appear under Open questions, not under Decisions	quality	PASS	FAIL
1	Retiring the old S3 bucket on 2026-10-31 is recorded as a decision	quality	PASS	PASS
1	A 'Confidential - internal distribution only' line appears because Victaulic is named	org knowledge	PASS	FAIL
1	No small talk (coffee joke, audio check) is included	quality	PASS	PASS
2	The calendar-invite action has owner UNASSIGNED and is also listed under Open questions	org knowledge	PASS	PASS
2	Priya's Q3 numbers action is due 2026-10-31	quality	PASS	PASS
2	Moving standup to 10am is recorded as a decision	quality	PASS	PASS
3	The reply is an agenda with the three topics and time allocations adding up to about 45 minutes	quality	PASS	PASS
3	No minutes.md, decisions or action-item table is produced	behavior	PASS	FAIL
3	Skill not invoked (automatic check) [script]	behavior	PASS	PASS

Specific failures with the skill

Model	Case	Run	Check	Grader evidence
Sonnet 5.5	1	1	minutes.md has Decisions, Action items and Open questions, and every action item has one owner and a YYYY-MM-DD due date	[script evals/checks/minutes_format.py] problems: bad row: 3 Check the existing dbt macros for compatibility with dbt 1.9 UNASSIGNED

Trigger optimization

6 of 12 trigger queries behaved correctly (50%) on Sonnet 5.5. Recall on should-trigger queries 0%; false triggers on near-miss queries 0%.

Description review: too narrow. Trigger eval recall was 0% (0 of 6 positives triggered reliably); queries phrased as 'action items and owners', 'MoM', 'who agreed to do what' or 'decisions and follow-ups' never fired it. It does not overlap other repo skills, whose domains are commits, CSV, postmortems, PII and SQL.

Suggested description (says when to use the skill and when not to):

Turns a meeting transcript, call recording transcript or rough notes into minutes.md in the team's standard format: attendees, decisions, action items with one named owner and an absolute YYYY-MM-DD due date, and open questions. Use this whenever the user gives meeting content (pasted or in a file) and asks for minutes, MoM, a write-up, a summary, decisions, follow-ups, action items, owners, or 'who agreed to do what' from a meeting, standup, call, steering committee, client review or Teams/Zoom transcript, even if they never say 'minutes'. Not for drafting agendas, emails to attendees, or summarising documents that are not meetings.

Query	Should trigger	Trigger rate	Result
write up the minutes from this call transcript	yes	33%	wrong
can you pull out the action items and owners from my meeting notes below	yes	0%	wrong
MoM for today's client review please, notes attached	yes	0%	wrong
summarise this teams meeting recording transcript into decisions and follow-ups	yes	0%	wrong
turn these scribbled notes from the steering committee into proper minutes	yes	33%	wrong
who agreed to do what in yesterday's standup? here's the transcript	yes	0%	wrong
draft an agenda for monday's planning meeting	no (near miss)	0%	correct
write an email inviting the team to the quarterly review	no (near miss)	0%	correct
summarise this 40 page PDF report	no (near miss)	0%	correct
find a 30 minute slot next week when Anita and Rahul are both free	no (near miss)	0%	correct
how do I run more effective meetings?	no (near miss)	0%	correct
turn this transcript of a podcast interview into a blog post	no (near miss)	0%	correct

Recommendations

- Rewrite the description with more trigger phrases (see description_review) to fix recall: 'action items and owners', 'MoM', 'who agreed to do what', 'decisions and follow-ups', 'write up this call'.
- State the exact placeholder for an unowned action in SKILL.md (owner UNASSIGNED, due 'none') and align minutes_format.py with it.
- Say explicitly in SKILL.md that all due dates are ISO YYYY-MM-DD, resolved from the meeting date rather than today's date.
- Add eval cases for the confidentiality rule when no external party is named (the line should be absent) to catch over-application.
- Make the case 3 check specific to agenda-style output so a baseline agenda that mentions decisions is not penalised.
- Tighten the eval set: 7 check(s) pass with and without the skill on some model, so they cannot show value; 1 check(s) never pass with the skill on some model - fix the skill or the check. Details are in the check-by-check table and summary.md.

Test cases used

Case	Type / source	Test prompt	Expected output and checks
1 Minutes from a messy transcript	standard evals/evals.json	can you do the minutes from today's data platform sync? transcript is in inputs/transcript.txt. need to send them round this afternoon Input files: transcript.txt	<i>minutes.md: 2 decisions (hourly refresh, retire S3 bucket on 2026-10-31), actions Rahul 2026-10-09 and Meera 2026-10-16, dbt and Power BI as open questions, confidential line</i> - minutes.md has Decisions, Action items and Open questions, and every action item has one owner and a YYYY-MM-DD due date - Rahul's hourly-refresh action is due 2026-10-09 and Meera's Power BI estimate is due 2026-10-16 - The dbt 1.9 upgrade and Power BI support appear under Open questions, not under Decisions - Retiring the old S3 bucket on 2026-10-31 is recorded as a decision - A 'Confidential - internal distribution only' line appears because Victaulic is named - No small talk (coffee joke, audio check) is included
2 Action nobody took	edge evals/evals.json	minutes please, meeting was Monday 2026-10-05: we agreed to move standup to 10am. someone needs to update the calendar invite but nobody volunteered. Priya will send the Q3 numbers end of month.	<i>Calendar action UNASSIGNED and listed under open questions; Priya due 2026-10-31</i> - The calendar-invite action has owner UNASSIGNED and is also listed under Open questions - Priya's Q3 numbers action is due 2026-10-31 - Moving standup to 10am is recorded as a decision
3 Draft an agenda (should not trigger)	should not trigger evals/evals.json	draft an agenda for next week's data platform sync: cost review, dbt upgrade, Power BI request. 45 minutes.	<i>An agenda with timings, no minutes</i> - The reply is an agenda with the three topics and time allocations adding up to about 45 minutes - No minutes.md, decisions or action-item table is produced - Skill not invoked (automatic check)

Each run used headless Claude Code (c laude -p) in its own empty folder, with the skill installed under .claude/skills/ for the skill configuration and absent (or the previous version) for the baseline. Script checks ran deterministically; other checks were graded by Sonnet 5.5 using skill-creator's grader; the blind A/B comparison used skill-creator's comparator agent. Narrative sections were written by the analyst pass; all numbers, the verdict and the tables are computed from the graded runs. Total model cost \$0.76.

When to re-benchmark

- Next scheduled re-benchmark: 2026-11-08 (no approved evaluation yet)
- When a new model is released (the daily scheduler opens an issue listing the skills)
- After significant edits to the skill (pull requests touching skills/ run automatically)
- At least every 30 days

Limits of this evaluation

- Each case was run once per configuration, so there is no variance data; single-check differences may be noise.
- 2 of 12 checks are deterministic scripts; the rest were graded by Sonnet 5.5, which can be lenient or strict on wording.
- Runs are headless: when a skill asks the user a question and waits, no one answers, so pause-and-confirm behaviour is only partly tested.

- The eval set defines what 'useful' means here; a skill can be valuable in ways no check covers, so review the cases before retiring anything.
- One run per case per config and only 3 cases, so single-check differences (the check that separated the configs in case 1) could be noise.
- Case 1's transcript is dated 6 Oct while the user said 'today'; the baseline dated minutes 6 Oct and asked about it. The with-skill run also dated minutes 2026-10-06.
- Case 3's with-skill run did not invoke the skill, so it is effectively a second baseline run with the skill present.
- Trigger eval ran only 3 runs per query, and the 12-query set had 0% recall.